

# Why our RBI response says what it says

Companion to the consultation response on RBI's Draft Guidance on Regulatory Principles for Model Risk Management, 2026. For each recommendation: exactly how it helps Krim, exactly who it is good for in the industry, whether it survives the disinterested-expert test — and the downsides and failure modes.

## THE POSTURE, IN ONE PARAGRAPH

We respond as a **supporter urging the regulator to hold the line**, not as a vendor asking for favours. We never ask RBI to require our product; we ask RBI to require the *properties* our product was built to have, and we volunteer to be bound by the same obligations we propose for others. A submission asking for less regulation would be ignored; one transparently selling KrimOS would be discarded. One that strengthens the regulator's own draft, paragraph by paragraph, gets read.

**The honest core of it:** our commercial interest lives less in the *content* of these recommendations than in their *selection and emphasis*. Most of what we say, a disinterested expert would also say. What a disinterested expert would not do is put pre-execution control at the centre of the letter. We do, because it is our thesis. That is what "declared interest" means, and we declare it.

## THE TEST WE APPLIED TO EVERY RECOMMENDATION

*Would a person with no product and no stake — an academic, a supervisor, a bank CRO — independently recommend this?* Where the answer is "yes," the recommendation is industry-first and Krim benefits as a side effect. Where the answer is "yes, but with less emphasis" or "only partly," the vendor fingerprint is real, and this document says so.

## THE CLASSIFICATION, AT A GLANCE

| # | RECOMMENDATION                 | HELPS KRIM?                                                              | GOOD FOR THE INDUSTRY?                                         | DISINTERESTED-EXPERT TEST                                                             |
|---|--------------------------------|--------------------------------------------------------------------------|----------------------------------------------------------------|---------------------------------------------------------------------------------------|
| 1 | Standardised vendor disclosure | Indirectly, strongly — disadvantages opaque vendors; we comply by design | Yes — fixes an information asymmetry that crushes small REs    | <b>Passes.</b> Classic market-failure remedy                                          |
| 2 | Prevention-capable controls    | Directly, maximally — our thesis written into regulation                 | Yes — irreversible actions need gates, not dashboards          | <b>Passes on content; the emphasis is ours</b>                                        |
| 3 | Decision-grade records         | Directly, strongly — native to our architecture, costly to retrofit      | Yes — disputes, examiners and consumer grievance all need them | <b>Passes with an accent</b> — "by construction" is our phrasing, chosen deliberately |
| 4 | Action-level override          | Modestly — Kupa is this                                                  | Yes — removes the operator's all-or-nothing dilemma            | <b>Passes cleanly</b>                                                                 |
| 5 | Sequenced implementation       | Only indirectly — costs us near-term                                     | Yes — RBI's own survey shows the cliff                         | <b>Passes</b> — industry would ask for it (minus our firm end date)                   |

|   |                        |                                                   |                                                                 |                                            |
|---|------------------------|---------------------------------------------------|-----------------------------------------------------------------|--------------------------------------------|
| 6 | Broad model definition | Directly — larger governed surface, larger market | <b>Split</b> — good for system and consumers; burdens small REs | <b>Contested</b> — the only one of the six |
|---|------------------------|---------------------------------------------------|-----------------------------------------------------------------|--------------------------------------------|

## RECOMMENDATION BY RECOMMENDATION

### 1 • Standardise the vendor disclosure the draft already requires

#### HOW IT HELPS KRIM

Closed foundation-model vendors and API resellers structurally cannot meet deep disclosure; KrimOS is validatable by design. A common schema converts our architecture from a differentiator we must argue for into a procurement requirement buyers must ask about. Volunteering to be bound costs us nothing and buys the whole letter credibility.

#### WHY IT IS GENUINELY GOOD FOR THE INDUSTRY

Paragraph 48 makes documentation a matter of bilateral negotiation. SBI can extract real artefacts from any vendor; a base-layer NBFC cannot. That is a textbook information asymmetry, and standardisation is the textbook remedy (the same logic as SOC 2 or the EBA's outsourcing registers). The direct beneficiaries are the *smallest* regulated entities — the ones with obligations under para 46 and no leverage to discharge them.

#### DISINTERESTED-EXPERT TEST

Passes. An academic in financial regulation would recommend this with no prompting.

#### DOWNSIDES / FAILURE MODES

- A weak schema becomes checkbox compliance and could even *help* opaque vendors ("we filled in the template").
- RBI may see schema-writing as beneath principles-level guidance and route it to an SRO, where it dies in committee.
- A standard weaker than our own disclosure practice levels the field against us. Net: asymmetric in our favour; worst case is being ignored.

## 2 • Prevention-capable controls for high-autonomy AI

### HOW IT HELPS KRIM

Maximally and directly: this is the founding thesis as regulatory language. If the final Guidance says controls for high-autonomy models must be able to stop a non-compliant action before execution, the market's compliance question changes from "do you have monitoring?" to "can you gate the action?" — the question KrimOS exists to answer. The single highest-value sentence we could get into the final text; internally we should be honest that we would pay to have it adopted.

### WHY IT IS GENUINELY GOOD FOR THE INDUSTRY

A lending action cannot be unmade; monitoring observes harm, only a gate prevents it. Action-gating before consequential effects is mainstream AI-safety architecture (EU AI Act Article 14's oversight-in-time; NIST AI RMF's govern-before-deploy posture) — and the draft itself already lists "verify and corroborate model outputs prior to their use" (para 54(1)(ii)); we ask RBI to promote its own clause. One subtlety keeps the ask honest: a manual maker-checker queue *is* prevention-capable. Banks satisfy this today at low volume; the requirement only forces automated gating at automation volume. Nobody is forced to buy anything — they are forced to *choose* between human-speed operations and a real gate.

### DISINTERESTED-EXPERT TEST

Passes on content, not on emphasis. A neutral expert would make this point once, mid-list. We built the letter's centrepiece out of it. That emphasis is the vendor fingerprint, mitigated by the "property, not a product" line — not eliminated.

### DOWNSIDES / FAILURE MODES

- **The backfire scenario, and it is real:** if prevention-capability makes high-autonomy AI look expensive, the rational bank response is to avoid high-autonomy deployments entirely — shrinking the near-term market for autonomous systems, including ours. We accept the risk because autonomy's economics are coming regardless; the question is whether it arrives regulated-and-permitted or informally tolerated until the first scandal. Krim wins in the first world, and so does the industry.
- Scope misreading: the draft governs *models*, our lens is *actions*. Scoped to "higher autonomy tiers" so it cannot be read as demanding gates on credit scorecards — still the recommendation most likely to be misread.
- Vague adoption ("entities should consider preventive controls") lets incumbent four-eyes processes claim compliance and changes nothing.

### 3 · Decision-grade records, produced by construction

#### HOW IT HELPS KRIM

Deterministic per-action records are what Krim-Nyāya produces natively; retrofitting them onto a conventional stack is expensive. The property is the product, same as recommendation 2.

#### WHY IT IS GENUINELY GOOD FOR THE INDUSTRY

Three constituencies need these records and currently don't get them: examiners (para 63 asks for review of interventions, overrides, incidents and near misses — impossible from reconstructed logs); REs themselves in disputes (reconstructed logs degrade exactly when litigation starts); and — the strongest link — **consumers**: para 25 requires grievance redressal for consumer-facing models, and a consumer cannot meaningfully challenge an AI-influenced decision unless a record exists of what the decision was checked against. Decision-grade records are the difference between grievance redressal as a right and as a gesture.

#### DISINTERESTED-EXPERT TEST

Passes with an accent. A neutral records expert would specify the *outcome* — contemporaneous, complete, tamper-evident, at action level. Our phrase "produced by construction" points at an *architecture* — ours. We chose the stronger phrasing deliberately; the neutral alternative existed, and this document records that we didn't use it.

#### DOWNSIDES / FAILURE MODES

- Genuinely burdensome for legacy incumbents; expect lobbying to soften it to "adequate records" (which changes nothing).
- "By construction" may read as architecture-prescription, which regulators resist.
- Storage and privacy interactions (per-action records × ten years × millions of actions) are real; if probed, the honest answer is that storage is cheap relative to disputes — have it ready.

### 4 · Action-level override, distinct from the kill switch

#### HOW IT HELPS KRIM

Modestly. Kupa, the command centre, is action-level supervision — but this is the least commercially loaded ask in the letter.

#### WHY IT IS GENUINELY GOOD FOR THE INDUSTRY

Without action-level intervention, an operator facing a misbehaving system has two buttons: "let it run" and "shut down everything." History across industries says nobody presses the second button until after the harm — the cost of a full stop is too visible, the cost of continuing too diffuse. Action-granularity gives supervisors a scalpel, which protects the bank's operators as much as its customers, and makes para 60's human-in-command real rather than ceremonial.

#### DISINTERESTED-EXPERT TEST

Passes cleanly. Control-engineering orthodoxy applied to AI operations.

#### DOWNSIDES / FAILURE MODES

- Risk of misreading as mandatory human review of *every* action — which would kill automation economics for everyone, us included. The text asks for the *capability* to intervene; summaries may blur that.
- Mild tooling-cost pushback expected.

## 5 • Sequenced implementation with a firm end date

### HOW IT HELPS KRIM

Indirectly at best — and in the near term it *costs* us: a slower forcing function means slower regulation-driven demand. The gains are credibility (the one recommendation that visibly runs against our interest) and a live market: our future customers survive the transition instead of freezing AI adoption or drowning in paper compliance. We deliberately spend advantage here to make recommendations 1–4 credible.

### WHY IT IS GENUINELY GOOD FOR THE INDUSTRY

RBI's own survey evidence ( $\approx 15\%$  interpretability tooling,  $\approx 18\%$  audit logs,  $\approx 21\%$  drift monitoring,  $\approx 1/3$  board oversight among AI-using lenders) says the sector cannot comply on day one. An impossible cliff produces fake compliance and blanket forbearance — worse for everyone than a sequenced ramp. Industry associations will ask for phasing themselves; the part they won't ask for is our firm end date, which is what keeps phasing from becoming indefinite delay.

### DISINTERESTED-EXPERT TEST

Passes. Any implementation economist would say the same.

### DOWNSIDES / FAILURE MODES

- "Proportionate to entity size" is exactly the clause lobbies stretch forever. Our hedge is the firm end date; **if RBI adopts the phasing without the end date, we will have argued for the indefinite postponement of our own market** — the single biggest strategic risk in the letter.
- The cynical reading (vendor recommends a timeline matching its own go-to-market) is true and unanswerable; the recommendation stands on RBI's own survey data anyway.

## 6 • Keep the broad definition of "model"

### HOW IT HELPS KRIM

Directly: the broad definition (para 7(3) — anything that takes inputs, applies logic, and materially shapes decisions, down to the pricing spreadsheet) means more of the lending stack needs governance — a larger addressable problem for a governance-layer company. Of the six, this is where our interest and the burden on parts of the industry pull hardest in opposite directions, so it deserves the most internal honesty.

### WHY IT IS GENUINELY GOOD — AND FOR WHOM

For the *system* and for *consumers*, unambiguously: perimeter-follows-function is the only anti-arbitrage definition, and a narrow one would be obsolete before publication (yesterday's spreadsheet is tomorrow's agent tool). For the *industry*, genuinely mixed: large banks can absorb the inventory burden; a base-layer NBFC discovering that a thousand spreadsheets are now regulated models cannot. A disinterested SME-lending advocate would oppose us here, and would not be crazy.

### DISINTERESTED-EXPERT TEST

Contested — the only one of the six. Systemic-risk experts side with us; small-lender advocates don't. Our answer to the tension is the *pairing*: recommendation 6 keeps the bar high, recommendation 5 spreads the runway. The two must survive together — the broad definition without sequencing is genuinely harsh, and we should not pretend otherwise internally.

### DOWNSIDES / FAILURE MODES

- Inventory-burden blowback lands on the constituency least able to complain effectively.
- "We anticipate pressure... urge RBI to resist" presumes other respondents' positions — mildly combative, accepted deliberately.

---

## WHAT WE CONSCIOUSLY DID NOT SAY

- **Nothing on data risk, privacy, or model-input governance** — real topics, not our lane; opining would dilute authority.
- **No comment on red-teaming standards (para 55) or customer AI-disclosure (para 59)** — supportable, but padding; six sharp recommendations beat nine soft ones.
- **No claims of deployments, customers, or India operations** — none exist to claim; the claim floor holds here especially.
- **No request for vendor recognition or certification** — commercially tempting, and the classic regulatory-capture ask; it would poison the well.
- **"Safe superintelligence" does not appear** — brand language that reads as science fiction to a model-risk supervisor. The letter says "safe autonomous systems for regulated financial work." Same substance, supervisory register.
- **We kept "produced by construction" over the neutral "contemporaneous and complete"** — a deliberate choice, recorded above.

**Primary-source verification — completed 6 July 2026.** The letter was checked line-by-line against the actual draft (16 pages, paragraphs 1–64), then independently gate-checked. Every characterisation and verbatim quote confirmed. Recommendation 1 was narrowed (paras 47–48 already contained half the original ask); recommendation 2 was re-anchored on para 54(1)(ii)'s own "prior to their use" clause; four citation-precision fixes applied (para 51 lists limiting usage as an example mitigant, not a mandate; para 54(2) permits "system-level controls or model design features"; the FREE-AI report is not this Guidance's formal parent — the draft anchors to Utkarsh 2029; paras 9–13 cover governance broadly, not the Board alone).

---

## PRE-SUBMISSION CHECKLIST

- ☒ Primary-source check — done 6 July 2026; recommendations 1 and 2 revised as a result.
- ☐ Confirm the FREE-AI survey figures (15/18/21/~1/3) against the FREE-AI Committee report itself.
- ☐ Confirm press-release number and dates against the published release (the draft PDF header uses placeholders).
- ☐ Fill: signatory name, title, email (personal or hello@, not sales@); legal entity line; date.
- ☐ Confirm the submission channel and format; send; retain the acknowledgement.
- ☐ Then: publish the response on the site, and LinkedIn post 1 goes live ("Ours goes in this week" must already be true).