


I N V E S T M E N T M E M O R A N D U M



KRIM

Sovereign Safe Superintelligence for Autonomous Banking Infrastructure

Powered by Kendra—the intelligence runtime where every AI action is validated before execution

Parameter	Details
Instrument	Compulsory Convertible Preference Shares (CCPS)
Raise Amount	\$1,500,000
Valuation (Pre-Money)	\$7.5M
Conversion	Converts at \$15–20M Seed round (target: 6–15 months)
Conversion Discount	20% discount to Seed price
Trigger	~\$2M ARR milestone
Entity	Krim AI Inc.
Jurisdiction	US (New York) + India operations

February 2026

Prepared by Shri Vishwa Nath Jha, Founder & CEO/CTO/CAIO

CONFIDENTIAL

Executive Summary

Metric	Detail
48M+ Borrowers Engaged (hist.)	Vishwa Nath Jha previously built and scaled India's largest voice AI platform for debt collection at Saarthi.ai (2017–2025)
200M+ Calls Processed (hist.)	Saarthi.ai engaged with 50+ banks and NBFCs with zero compliance violations
\$10M+ Active Pipeline	Yes Bank is in active onboarding; SBI is advancing toward a letter of intent—both before any institutional fundraise.
\$1M ARR Floor	Based on Yes Bank's active engagement alone (₹70L MRR, bottom-up by DPD segment); SBI and expansion modules represent pure upside
Working GTM Product	Currently in testing, demo, and iteration; production launch targeted April/May 2026
2 Co-Founders	CEO/CTO/CAIO (architecture + domain), COO (operations + US market + enterprise sales)

Krim is building KrimOS—Sovereign Safe Superintelligence for Autonomous Banking Infrastructure—powered by Kendra, the intelligence runtime that sits between foundation models and enterprise applications, where every AI action is validated before execution. Over the next 3–5 years, AI maturity will make human middleware in banking operations obsolete. The question is not whether banking becomes autonomous, but whether autonomous banking can be made safe. Krim answers that question.

While every major AI lab invests in post-generation safety (guardrails, filters, RLHF), none can answer the foundational question that autonomous banking demands: how does the system know that what it knows is true? In consumer applications, this gap is tolerable. In banking—where every AI action carries legal, financial, and human consequences—it is existential. A 95% compliance rate at scale means hundreds of thousands of regulatory exposures per year. Current SaaS banking platforms digitise human workflows; Krim replaces them with autonomous operations where compliance violations are structurally impossible.

We call this category Epistemological AI—the first AI architecture where systems validate their epistemic warrant before any action executes. Krim's framework provides 26 structurally distinct pre-execution validators, rooted in Navya-Nyāya, a 14th-century Indian formal logic system—the most rigorous pre-modern epistemology for validating knowledge claims—now translated into computational form for the first time.

The investment opportunity: \$1.5M CCPS at a \$7.5M pre-money valuation, converting at a 20% discount into an expected \$15–20M Seed round within 6–15 months. The founder's prior pedigree—48M borrowers, 50+ banks, zero violations—has already produced active bank onboarding before any institutional capital.

1. The Investment Thesis: The First Epistemological AI Company

Banking is about to undergo the same transformation that manufacturing underwent with robotics: human middleware—the thousands of people manually executing operational workflows—will be replaced by autonomous AI systems. The current generation of banking SaaS platforms (Credgenics, DPDzero, Yubi) digitises human processes. The next generation will replace them entirely. Krim is building the safe superintelligence infrastructure that makes this transition possible.

Every AI system deployed in banking today has a structural defect: it cannot prove its reasoning is valid before acting. Krim is building the only architecture that solves this—Kendra, the intelligence runtime that validates reasoning before execution, reducing compliance violations by an order of magnitude and making the most dangerous failure modes structurally impossible.

The AI industry has converged on a single safety architecture: generate first, filter second. Guardrails, content filters, LLM-as-judge, RLHF—all operate after generation. This creates a hard ceiling: post-hoc detection catches 60–80% of novel errors in production. The remaining 20–40% require a fundamentally different approach.

In consumer applications, this ceiling is tolerable. In regulated industries—where every AI action carries legal, financial, and human consequences—it is existential. A 95% compliance rate at scale means hundreds of thousands of regulatory exposures per year.

Krim is pioneering a new category: Epistemological AI—systems that validate the reasoning warrant for every decision before execution. Unlike guardrails that filter outputs or alignment techniques that shape model behavior, Epistemological AI asks a fundamentally different question: does this system have valid grounds for knowing what it claims to know? The framework draws on Navya-Nyāya, a 14th-century Indian formal logic system that provides the most rigorous pre-modern epistemology for validating knowledge claims—now translated into computational form for the first time.

Why This Matters Now

Force	What Changed	Implication
AI Economics	Cognitive cost fell 100× in 24 months	AI judgment now cheaper than human middleware
Agentic AI	Autonomous agents shipping to production	Every agent action is a potential regulatory event
Regulatory Pressure	CFPB enforcement up 340%; RBI tightening with DPDPA	\$15B+ compliance modernisation spend being forced

2. The Problem: AI’s Epistemological Crisis

Every major AI lab talks about safety. None can answer: how does your system know that what it knows is true?

OpenAI builds guardrails. Anthropic builds constitutions. Google builds filters. Each shares the same assumption: the right response to unreliable AI is to catch bad outputs after generation. This is end-of-line quality inspection—a practice manufacturers abandoned decades ago.

The Post-Disbursal Credit Operations Problem

Metric	Scale	Problem
Global credit market	\$140B+ in post-disbursal ops	85% still manual processes
US compliance spend	\$50B annually	Fines keep climbing despite spending
India credit growth	\$3T+ in new credit by 2030	Infrastructure doesn’t exist for this scale
Delinquency rates	15-year highs across markets	Manual operations breaking under volume
Agent turnover	40–60% annually	Institutional knowledge vanishes every 18 months

The Trust Gap isn’t technical. AI is capable enough. The gap is epistemic—organisations cannot verify that their AI knows what it’s doing, as opposed to statistically guessing right most of the time.

The Validation Crisis: Ten Cognitive Defects Documented at Enterprise Scale

Scaling voice AI across 50+ regulated financial institutions—processing hundreds of millions of interactions—revealed ten classes of failures that existing architectures cannot solve. These aren’t theoretical—they’re documented from production deployments at enterprise scale:

#	Cognitive Defect	Production Evidence	Why Post-Hoc Detection Fails
1	Hallucination Under Constraint	0.3% of calls contained compliant-sounding responses with disguised threats	Output appears compliant; only epistemic analysis catches the warrant violation
2	Emergent Non-Compliance	Optimising for “help customer” vs “maximise collection” increases violations 4× despite identical rules	Same guardrails, opposite outcomes—rules-based systems can’t detect goal misalignment
3	Context Window Amnesia	12% of calls exceeding 20 minutes showed degraded constraint adherence	Compliance erodes gradually; no single output triggers a filter
4	Multi-Step Reasoning Failures	8% of violations were compositional—each step valid, but the workflow invalid	Individual action passes; only sequence-level validation catches the failure

#	Cognitive Defect	Production Evidence	Why Post-Hoc Detection Fails
5	Adversarial Brittleness	Red team achieved 15% success rate inducing violations through strategic prompting	Prompt injection bypasses output filters by manipulating reasoning, not outputs
6	Statistical Drift	Compliance degraded from 99.8% to 97.2% over 3 months as models learned aggressive tactics increase payment rates	Slow drift is invisible to threshold-based monitoring until damage accumulates
7	Composition Failures	5% of violations occurred at multi-agent handoff boundaries—agents individually compliant, handoffs non-compliant	Voice agent: “follow up tomorrow.” SMS agent: sends 2 hours later (TCPA requires 24-hour minimum). As systems scale to 8–20+ agent types, handoff complexity grows quadratically
8	Temporal Violations	Actions compliant at validation become violations at execution due to time-dependent regulatory windows	Static validation at generation time misses dynamic temporal constraints
9	Implicit Bias Leakage	12% higher denial rates in protected geographies despite identical credit profiles	Bias is statistical, not textual—content filters don’t detect distributional skew
10	Validation Bypass	30–40% of keyword filters bypassed through rephrasing	Keyword-based detection is structurally defeatable; semantic-level validation is not

These ten defects are not edge cases—they are structural properties of post-hoc safety architectures. No amount of better guardrails, stronger filters, or smarter judges solves them, because the failure occurs before the output is generated—in the reasoning warrant itself. This is why Krim builds pre-execution validation: to prevent the defect at the cognitive layer, not detect it at the output layer.

3. The Solution: Epistemological AI

Epistemological AI inverts the safety paradigm. Instead of Input → Generate → Filter → Output, the architecture becomes Input → Validate Epistemic Warrant → Generate with Proven Grounds → Output.

Every action must demonstrate valid grounds for knowing before execution. If the warrant doesn't hold, the action doesn't execute—not because a filter caught it, but because the architecture won't permit unjustified inference.

KrimOS and Kendra: The Safe Execution Runtime

KrimOS is a five-layer platform: Kriya (atomic validated primitives), Karta (autonomous AI coworkers), Kendra (the intelligence runtime), Kula (personalised digital twins for bank staff), and Kupa (command centre for managers).

Kendra is the core runtime layer—containing seven subcomponents: Krim-Core (lifecycle orchestration), Krim-Govern (7-level policy hierarchy), Krim-Nyāya (26 pre-execution epistemological validators), Krim-Fabric (knowledge graph and cross-system context), Krim-Ledger (immutable audit trail with cryptographic hash chains), Krim-Learn (federated learning and cross-tenant intelligence), and Krim-Sense (telemetry and observability). Compliance is encoded as structural constraint within Kendra. The most dangerous failure modes—calling the wrong borrower, acting on closed accounts, violating contact-hour regulations, fabricating data—are made structurally impossible through Kendra's pre-execution validation.

Layer	Name	Function
1. Atomic Primitives	Kriya	Indivisible units of validated work—each with built-in compliance constraints
2. AI Co-workers	Karta	Autonomous agents composed from Kriya primitives; each action epistemically validated
3. Intelligence Runtime	Kendra	Orchestration layer with pre-execution validation as first-class infrastructure
4. Digital Twins	Kula	Digital witnesses for human staff—compose, spawn, and govern Karta workers
5. Command Centre	Kupa	Enterprise-wide visibility, reporting, and human-in-loop governance

Kendra's pre-execution validation is what separates KrimOS from every other platform in the market. Other failure modes—beyond the structurally impossible ones—are reduced by an order of magnitude compared to post-hoc filtering.

Kula Digital Twins: How Banks Actually Deploy KrimOS

A Kula is a digital twin of a human role. It encodes the judgment, domain knowledge, and decision boundaries of a specific operational function—then governs a fleet of Karta AI agents that execute autonomously within those boundaries. Banks don’t deploy “AI agents.” They deploy digital twins of their best people.

Kula Digital Twin	Human Role It Mirrors	What It Automates	Karta Agents It Governs
Collections Strategist	Senior collections manager	Portfolio segmentation, risk bucketing, strategy assignment. Dynamically adjusts strategy based on borrower response patterns.	Outbound caller agents, negotiation agents, escalation agents
Voice Operations Lead	Call centre team lead	Call scheduling within legal windows, script selection, real-time compliance monitoring during live calls, post-call dispositioning.	Voice agents (outbound/inbound), IVR agents, callback agents
Settlement Negotiator	Senior recovery officer	Autonomous settlement offer generation within pre-approved parameters. Validates offer against account balance, loan terms, regulatory limits.	Negotiation agents, payment plan agents, approval routing agents
Field Operations Controller	Field ops manager	Field visit scheduling, agent route optimisation, visit outcome tracking, photo/document capture verification.	Field coordination agents, document verification agents, geo-routing agents
Compliance Auditor	Internal audit / compliance officer	Continuous monitoring of all agent actions against 26 validators. Generates audit trails, flags anomalies, produces regulatory reports.	Audit trail agents, anomaly detection agents, report generation agents
Back-Office Processor	Operations analyst	Payment reconciliation, account status updates, dispute processing, refund initiation, loan restructuring paperwork.	Reconciliation agents, status update agents, document processing agents

Each Kula digital twin is created by a human expert—encoding their decision logic, compliance boundaries, and escalation criteria. The human doesn’t disappear. The human’s judgment scales to thousands of simultaneous operations while maintaining the quality, compliance awareness, and contextual sensitivity of the original expert.

The 26 Pre-Execution Validators

The AI industry calls every incorrect output a “hallucination.” It is a catch-all that explains nothing. Krim’s epistemic framework provides 26 structurally distinct validators—6 knowledge source checks, 13 error type checks, and 7 semantic compatibility checks—each targeting a different failure mode with a different cause, detection method, and prevention mechanism.

I. 6 Knowledge Source Validators (Pramāṇa)

Before any AI action, every input is classified by source type and assigned an epistemic weight. Higher-ranked sources receive greater trust. Lower-ranked sources trigger additional verification.

#	Validator	Epistemic Weight	What It Checks	Example
1	Direct Perception (Pratyakṣa)	Highest	Is this data from a verified operational source?	Account balance from CBS = highest trust
2	Inference (Anumāna)	High	Is this computed from validated premises using valid logic?	Predicted payment probability from verified inputs
3	Analogical Comparison (Upamāna)	Medium	Is this similar-case match from a structurally comparable scenario?	Similar borrower profile resolved via Plan X
4	Authoritative Testimony (Śabda)	Source-dependent	Is this from a trusted authority?	RBI circular = authoritative; unverified CRM note = low
5	Necessary Presumption (Arthāpatti)	Conditional	Is this the only logically necessary conclusion?	No payment + no dispute = presumed default intent
6	Validated Absence (Anupalabdhi)	Blocking	Is the absence of expected data itself a signal?	No phone on file = blocks call attempt automatically

II. 13 Error Type Validators (Doṣa)

Each error type maps to a specific AI failure mode with distinct detection and prevention logic. Current AI safety treats all of these as “hallucination.” Krim diagnoses and prevents each separately.

#	Error Validator	What It Catches	AI Failure Prevented	Severity
1	Misidentification (Anyathā-khyāti)	Wrong entity selected	Calling wrong borrower	CRITICAL
2	Fabricated Entity (Asat-khyāti)	Reference to non-existent data	Citing non-existent account	CRITICAL
3	Inverted Logic (Viparīta-khyāti)	Conclusion contradicts premises	Offering hardship plan to fully paid loan	HIGH
4	Unresolved Ambiguity (Saṁśaya)	Acting despite insufficient clarity	Proceeding when unclear which account	MEDIUM

#	Error Validator	What It Catches	AI Failure Prevented	Severity
5	Fallacious Reasoning (Hetvābhāsa)	Invalid logical structure	Treating correlation as causation	HIGH
6	Unverified Premise (Asiddha)	Acting on unconfirmed assumption	Using unverified borrower status	HIGH
7	Self-Contradiction (Viruddha)	Output contradicts own inputs	Promising waiver while demanding full payment	CRITICAL
8	Ignored Counter-Evidence (Satpratipakṣa)	Contradictory data not considered	Ignoring payment receipt while pursuing collection	MEDIUM
9	Temporal Violation (Kāla-doṣa)	Time-based rule breach	Calling outside legal hours	CRITICAL
10	Pre-Existence Error (Prāgabhāva)	Acting on entity not yet created	Processing future-dated account	CRITICAL
11	Post-Destruction Error (Pradhvaṃsābhāva)	Acting on terminated entity	Collecting from deceased borrower	CRITICAL
12	Absolute Non-Existence (Atyantābhāva)	Pure hallucination	Acting on entity that never existed	CRITICAL
13	Identity Confusion (Anyonyābhāva)	Two entities swapped	Applying Person A's terms to Person B	CRITICAL

III. 7 Semantic Compatibility Validators (Yogyatā)

Even with valid sources and no reasoning errors, an action can be semantically absurd. All seven compatibility checks must pass before any action proceeds.

#	Compatibility Check	What It Validates	Example Failure Blocked
1	Temporal Compatibility (Kāla)	Are dates, times, and windows valid?	Scheduling payment for Feb 30; calling at 3 AM
2	Numeric Compatibility (Saṅkhyā)	Are amounts and quantities sensible?	Reducing balance to -\$500; EMI of \$0
3	State Compatibility (Avasthā)	Is action compatible with entity's current state?	Calling deceased borrower
4	Product Compatibility (Utpatti)	Does action match the product type?	Offering 0% APR on fixed-rate mortgage
5	Channel Compatibility (Mādhyaṃa)	Is the communication channel appropriate?	Sending SMS to landline
6	Status Compatibility (Sthiti)	Is the action valid given lifecycle status?	Processing ACH to closed account
7	Regulatory Compatibility (Niyama)	Does it comply with applicable regulatory context?	Applying wrong state law

6 Knowledge Source Validators + 13 Error Type Validators + 7 Semantic Compatibility Validators = 26 structurally distinct pre-execution checks. Each validator targets a specific failure mode with a specific cause, detection method, and prevention mechanism. Current AI safety has one word for all 26 failure modes: “hallucination.” This is why post-hoc safety has a ceiling—you cannot fix what you cannot diagnose.

Intellectual Property Strategy

The 26-validator computational framework represents a first-of-kind formalisation of Navya-Nyāya epistemology into executable AI safety infrastructure. Krim’s IP strategy includes: (a) provisional patent filings covering the computational translation of epistemic validators into pre-execution validation logic, (b) trade secret protection for the specific implementation within KrimOS, and (c) academic documentation through a three-paper series in collaboration with Prof. Girish Nath Jha at JNU, establishing priority of invention and publication record.

4. Why Voice AI Cannot Scale Without Valid Reasoning

Every platform shift in enterprise AI follows the same arc: a technology wave creates massive adoption, then a trust problem emerges that demands a new category of infrastructure.

Era	Wave	Market Created	Trust Problem	What Solved It
2016–2020	Conversational AI	\$14B+ market by 2025	Conversations were robotic, scripted	Natural language models (GPT, BERT)
2020–2024	Voice AI	\$45B+ market by 2025	Voice could talk but couldn't think—no compliance awareness	LLM-powered voice agents
2024→	Epistemological AI	Next category: validated intelligence	AI can reason—but cannot prove reasoning is valid before acting	Pre-execution epistemic validation (Krim)

The Scaling Paradox: More Capable Voice = More Dangerous Unvalidated Reasoning

Conversational AI solved reach—machines could talk to customers at scale. Voice AI solved naturalness—the conversations became human-like. But neither solved trust—the ability to guarantee that what the AI says and does is epistemically valid.

Voice AI without validated reasoning creates the worst possible outcome: a convincing AI agent that confidently makes wrong decisions. A chatbot hallucination is embarrassing. A collections agent hallucination is a CFPB enforcement action. An autonomous agent that offers an unauthorised settlement, calls outside legal hours, or misidentifies a borrower doesn't just fail—it creates regulatory liability at machine speed.

Sierra AI's \$10B valuation proves voice AI demand at scale. But even Sierra—with \$635M in funding, \$100M ARR, and a “constellation” supervisor model—suffered a coordinated jailbreak in December 2025 that exposed guardrail limits across a dozen customer agents. GigaML's 90%+ resolution rate with DoorDash is impressive for food delivery. In financial services, the 10% failure rate would be catastrophic.

What the Frontier Labs Are Doing—And What They're Not

Every major AI lab is investing in safety—but all are working within the same paradigm: making models reason better, then filtering outputs after generation. DeepSeek R1's chain-of-thought reasoning, Ai2's OLMo 3 transparency, Moonshot's Kimi K2 thinking modes, Anthropic's Constitutional AI—all make models smarter, more transparent, and more aligned. None of them validate whether a specific action in a specific regulatory context is epistemically justified before execution.

The Emerging Validation Landscape: Infrastructure Players Are Racing to This Problem

The market is validating this problem urgently. In December 2025 alone, major infrastructure players shipped production validation solutions—confirming that post-hoc detection has hit its ceiling:

Initiative	What It Does	Structural Limitation
vLLM HaluGate (Dec 14, 2025)	Token-level hallucination detection with 76–162ms latency overhead. Detects when models generate responses contradicting provided context.	Post-hoc detection after generation—catches errors but cannot prevent them structurally.
vLLM Semantic Router x AMD (Dec 16, 2025)	“Intelligence Control Plane” for AI governance with three lifelines: World Output (action checkpoints), World Input (border inspection), Long-Term State (memory governance). “Transform alignment from training-time aspiration into runtime institution.”	Focused on model routing and post-hoc guardrails, not pre-execution validation at the cognitive layer.

The pattern: infrastructure companies are racing to build validation layers because post-hoc detection fails in production. But they’re approaching from inference optimisation or model routing—not from first-principles epistemology. While others ask “did the model hallucinate?” KrimOS asks “is this action structurally valid to execute?” This is the difference between catching a bug in QA versus type-safety preventing the bug at compile-time.

The frontier labs are solving the model intelligence problem. Krim is solving the action validation problem—the prerequisite for autonomous banking. These are complementary—and Kendra, Krim’s intelligence runtime, is model-agnostic by design, benefiting from every advance the labs make. As models become more capable, the case for autonomous banking strengthens—and the need for validated reasoning before execution becomes non-negotiable.

5. Where Value Is Shifting: The Missing Layer in AI Infrastructure

At Davos 2026, Jensen Huang described AI as a “five-layer cake”—Energy → Chips → Cloud Infrastructure → AI Models → Applications. But his framework reveals a critical gap between Layer 4 (Models) and Layer 5 (Applications)—the Intelligence Runtime—that determines whether models can be deployed safely in regulated environments. This is where Krim operates: Kendra, the intelligence runtime for autonomous banking, sits at Layer 4.5 inside KrimOS. Current SaaS platforms at Layer 5 digitise human workflows. Kendra at Layer 4.5 makes those workflows autonomous—and safe.

Layer	Huang’s Stack	Current Leaders	Krim’s Position
5	Applications	Vertical AI companies, enterprise adopters	Krim’s credit operations applications sit here
4.5	Intelligence Runtime	No one	KrimOS: 26 pre-execution validators
4.2	Model Routing & Governance	AMD × vLLM SR, Envoy AI Gateway	Routes to right model; Krim validates right action
4	AI Models	OpenAI, Anthropic, Google, Meta, DeepSeek	Model-agnostic; consumes any foundation model
3	Cloud Infrastructure	AWS, Azure, GCP, Oracle	Deploys on standard cloud
2	Chips & Computing	NVIDIA, AMD, custom silicon	Hardware-agnostic
1	Energy	Utilities, sovereign infrastructure	—

Why Layer 4.5 Is the Highest-Value Position

Models are commoditising. Enterprise AI infrastructure spend reached \$18B in 2025, but performance gaps between frontier models are narrowing. DeepSeek, Qwen, LLaMA, and Mistral provide high-performance open-source alternatives. IBM’s Chief Architect declared: “We’re going to hit a commodity point.”

When models become interchangeable, the competitive moat shifts to what sits between them and the enterprise. Four infrastructure categories are now foundational: agent orchestration, model routing, safety and governance, and fine-tuning. Krim unifies all four into Kendra—a single intelligence runtime purpose-built for regulated industries.

The analogy: HTTPS didn’t make the internet faster. It made the internet deployable for commerce and banking. KrimOS—powered by Kendra, the intelligence runtime for autonomous banking—doesn’t make models smarter. It makes autonomous banking possible by ensuring every AI action is validated before execution.

Compounding Moat: Why This Position Deepens Over Time

Infrastructure layer plays are notoriously difficult—API gateways commoditised, model routers went open-source, guardrails became table stakes. Krim’s position is structurally different because the moat compounds across four dimensions that are individually hard and collectively rare:

1. Domain embedding depth. Every bank deployment embeds Krim deeper into credit operations workflows. The 26 validators are calibrated through production data across different regulatory regimes, borrower populations, and product types. Each deployment makes the next faster and the validator framework more precise. Competitors entering later face not just a technology gap but a calibration gap that takes years to close.

2. Switching cost architecture. KrimOS sits between the bank’s core systems and its AI agents. Once Kula digital twins encode an institution’s decision logic, compliance boundaries, and escalation criteria, replacing KrimOS means rebuilding the entire operational intelligence layer—not just swapping a vendor.

3. Interdisciplinary rarity. Building KrimOS requires simultaneous expertise in classical Indian epistemology, AI/ML engineering, financial services regulation, and production-scale voice AI. This combination exists nowhere else—not at OpenAI, not at Sierra, not at any collections SaaS company. A competitor would need to assemble a team that spans philosophy, engineering, and regulatory compliance—and then spend years in production before matching Krim’s validator calibration.

4. Network effects without data sharing. As more banks deploy KrimOS, the Karta agent library grows. Each customer’s new AI co-workers—purpose-built for specific operational workflows—become available to new customers on Day 1. The library compounds: launch with 10 agents, Year 1: 67, Year 2: 203+. New entrants deploy into a cold library. Krim customers deploy into a warm one.

These four dimensions compound multiplicatively. A competitor entering 18–24 months later faces not just a head start but a structural gap that widens with every deployment.

6. Competitive Landscape

The following analysis evaluates competitors across both industry-standard operational metrics and Krim’s proprietary epistemic validation framework.

Category 1: Voice AI Companies—Global / Scaled Players

Company	Funding / Scale	Strategy	Structural Limitation
Sierra	\$635M raised; \$10B valuation; \$100M ARR in 7 quarters. Bret Taylor + Clay Bavor.	Enterprise CX AI agents. “Constellation” supervisor model. Outcome-based pricing.	Horizontal CX—not domain-specific to credit ops. Supervisor is still post-hoc. Dec 2025 jailbreak exposed guardrail limits. No pre-execution framework.
GigaML (Giga)	\$61M Series A (Nov 2025); Redpoint, YC (\$23). IIT Kharagpur founders.	Voice AI for enterprise ops. DoorDash, Zepto. 90%+ resolution rate. Expanding into BFSI.	Ops automation, not intelligence runtime. 90% = 10% failure rate—catastrophic in regulated industries. Entering finance without compliance architecture.
Skit.ai	\$47.6M raised; ~\$22M rev (FY24); 500 employees	AI voice agents for US debt collection. 53K+ creditors. \$1B+ in consumer debt resolved.	Voice-only. No back-office, no reconciliation. Post-hoc compliance. No pre-execution validation.
Salient	\$60.5M raised; a16z + Matrix	AI phone agents for financial services. YC-backed.	Early-stage. Voice-first, no operations layer. No pre-execution validation.
Observe.AI	\$214M raised; \$100–250M rev; SoftBank-backed; 400+ employees	Contact centre intelligence + VoiceAI agents. QA, coaching.	Broad horizontal play. Observes—does not orchestrate. Post-hoc analysis.

Category 1B: India Voice AI—Recent Fundraises (2024–2026)

Over \$200M raised by Indian voice AI startups in 18 months, validating massive market demand. Every one operates at Layer 5 (applications)—none at Layer 4.5 (intelligence runtime). They are future Krim customers, not competitors.

Company	Funding	Strategy	Structural Limitation
Sarvam AI	\$53M Series A	Full-stack Indic AI: sovereign LLM under IndiaAI Mission	Foundation model play, not domain-specific. Pivoted from voice to sovereign LLM.
Nurix AI	\$27.5M Seed+Series A; \$33M valuation	Custom enterprise AI agents via NuPlay platform. Sub-500ms latency.	Horizontal play—no collections domain. No epistemic validation.
Mihup	\$16.3M+ total; IPO planned in 2 years	Proprietary ASR/NLP/TTS stack. 100K+ devices deployed.	Contact centre analytics. Observation layer, not intelligence runtime.

Company	Funding	Strategy	Structural Limitation
GreyLabs AI	\$11.2M total (\$9.6M Series A)	Agentic voice AI for BFSI. 50+ FI clients.	BFSI voice + analytics. No back-office, no reconciliation. Post-hoc QA.
Bolna	\$6.9M Seed (Jan 2026); General Catalyst, YC	Self-serve voice AI. 200K+ daily calls. 1,050+ paying customers.	Horizontal voice. No domain compliance layer.
Smallest AI	\$8.3M Seed (Aug 2025)	TTS (Lightning model), voice agents. Developer-first.	TTS/voice infra play. No domain logic or compliance architecture.

Category 2: Collections SaaS & AI Platforms

Company	Funding / Scale	Strategy	Structural Limitation
Credgenics	\$80M raised; ₹220Cr rev (FY25); \$340M valuation; 474 employees	India's #1 collections SaaS. 150+ FI clients. 98M loan accounts.	SaaS workflow tool, not intelligence runtime. Digitises human processes.
Spocto X (Yubi)	Acquired ~\$58M (2022); launched YuVerse.ai; acquired CreditNirvana (Mar 2025)	AI-driven collections for PSBs. Now building own agentic AI layer via YuVerse.	Part of Yubi ecosystem. Building competing platform—validates market but innovation constrained by acquirer roadmap.
DPDzero	\$10.8M total (\$7M Series A); Revenue 6x in 18 months	AI-powered collections for fintechs/NBFCs. 1.8M borrowers/month.	Collections-as-a-service, not intelligence runtime. No epistemic validation.
Rezolv	\$3.5M Seed (Mar 2025); Kissht co-founders	Full-lifecycle AI collections platform.	Collections SaaS—digitises existing process. Workflow tool, not reasoning layer.

Category 3: Guardrails, AI Safety & Routing Infrastructure

Company	Approach	Structural Limitation
AMD × vLLM Semantic Router	Open-source Intelligence Control Plane. AMD GPU-accelerated semantic routing. NeurIPS 2025 paper.	Routes requests to the right model—does not validate per-action reasoning. Post-generation hallucination detection.
Guardrails AI	Open-source LLM output validation against rules.	Post-hoc by design. Cannot prevent errors it hasn't anticipated.
Lakera	AI security: prompt injection, content moderation, PII.	Security layer, not reasoning layer. Catches attacks, not invalid inferences.
Constitutional AI	Training-time alignment via principles (Anthropic).	Model-level, not deployment-level. Doesn't validate per-action regulatory compliance.

The Capability Map

Capability	Voice AI	Collections SaaS	Guardrails / Routing	Krim
Voice automation	✓	X	X	✓
Back-office operations	X	Partial	X	✓
Multi-agent orchestration	X	X	Model routing only	✓
Compliance audit trail	Basic logs	Partial	Alert logs	Full epistemic audit
Hallucination detection	X	X	Post-generation	Pre-execution (26 types)
Pre-execution validation	X	X	X	✓ (26 validators)
End-to-end credit ops	X	Partial	X	✓

\$1B+ raised by voice AI companies targeting BFSI/collections in 18 months. All operate at Layer 5—digitising existing workflows. None at Layer 4.5—enabling autonomous operations. AMD × vLLM routes to the right model; Kendra, Krim’s intelligence runtime, validates the right action. They are complementary layers. The structural difference: Layer 5 companies improve the current operating model. Krim replaces it.

7. Founder-Market Fit: Why This Team

Yes Bank and SBI engaged with Krim based on the founder's track record, representing \$10M+ in active pipeline. Banks follow proven operational results—the founder's track record of 200M+ calls with zero violations across 50+ banks is why these engagements exist.

Shri Vishwa Nath Jha — Founder & CEO/CTO/CAIO

Founder & CEO of Saarthi.ai (2017–2025): built India's largest voice AI platform for debt collection from scratch. Personally architected the voice AI infrastructure that processed 200M+ calls across 50+ banks/NBFCs with zero compliance violations. Exited Saarthi.ai in a strategic founder transition—a deliberate decision to build Krim as a clean-sheet architecture. Epistemological AI requires fundamentally different system design; it could not have been retrofitted onto Saarthi's existing codebase.

There is no IP overlap, no operational entanglements, and no restrictions that prevent Krim from building or selling KrimOS. The founder's prior company has pending financial obligations (TDS, vendor, and PF liabilities) that are in active resolution—these are documented transparently in the Investor Protection Framework, including the temporary nominee directorship structure and its resolution path. The proof: India's largest banks are actively engaging with Krim—Yes Bank is onboarding and SBI is advancing toward a letter of intent. Investor representatives are welcome to verify engagement status directly with bank counterparts. Deep expertise in classical Indian epistemology; academic collaboration with Prof. Girish Nath Jha at JNU on Navya-Nyāya applications to AI safety.

Mr. Louis Oliphant Parkinson — Co-Founder & COO

Former CRO of Saarthi.ai. BSc Natural Sciences (Maths & Philosophy), Durham. Led C-suite sales and voice AI transformation at India's largest banks. Deep expertise in enterprise operations, P&L, legal, HR, and scaling human-AI hybrid teams. To be based in New York, bringing US market access for Seed-stage expansion.

Engineering Scaling Plan

Current state: The founder's triple role (CEO/CTO/CAIO) is a strength at this stage—zero translation loss between vision, architecture, and execution. He personally built the working GTM product now in bank testing and demo. Scaling plan: A senior engineering hire (VP Engineering or Principal Engineer) is budgeted within 3–4 months of close. This hire absorbs day-to-day engineering execution, enabling the founder to focus on architecture, epistemic framework development, and enterprise relationships. Target team size by Month 9: 6–8 engineers.

8. Traction, Product Status & Revenue Projections

Current Product Status

KrimOS is a working GTM product currently in active testing, demonstration, and iteration with bank teams. The system has been demoed to banking stakeholders and is being refined based on feedback from Yes Bank's technical and operations teams, with empanelment in final stages. Production launch is targeted for April/May 2026.

The path to full-stack intelligence runtime (Kendra) for autonomous credit operations requires approximately 3 additional months of intensive development—and is a continuous process shaped by bank teams' feedback. This iterative approach is deliberate: building in production with real bank requirements ensures the platform solves actual operational problems rather than theoretical ones.

Current Pipeline

Partner	Status	Revenue Basis	Significance
Yes Bank	Onboarding active; testing & iteration underway	\$1M ARR (₹70L MRR bottom-up; revenue floor)	Production onboarding started. Revenue from early months post-launch.
State Bank of India	LOI ready; 2-week trial to follow	Significant expansion potential	500M+ customers; India's largest bank. LOI expected before trial commencement.
HDB Financial Services	Discovery	Additive to base pipeline	HDFC Bank subsidiary
L&T Finance	Discovery	Additive to base pipeline	Major diversified NBFC

The \$1M ARR projection is based strictly on Yes Bank's active engagement alone—₹70L MRR from 100K borrowers, modeled bottom-up by DPD segment at 160 KWU per borrower (see Sales & Commercial Overview for detailed segment-level consumption model). SBI pilot conversion, HDB, L&T, and the founder's network across 50+ banks represent pure upside above this floor. Both Yes Bank and SBI engagements were secured through the founder's prior pedigree—before any institutional fundraise.

Revenue Projections (Conservative)

These projections reflect realistic Indian bank procurement timelines. Yes Bank onboarding is active; first revenue at ₹70L MRR (~\$1M ARR) is projected from Month 1 post go-live as KrimOS processes 100K borrowers across Pre-Due to DPD 90.

Timeline	ARR	Drivers	Milestone
Month 1 post go-live	\$1M	Yes Bank production goes live	First revenue (₹70L MRR)
Month 6	\$2.2M	Yes Bank borrower expansion (100K→220K); SBI trial underway	Seed trigger achieved
Month 9–12	\$3.5M	Yes Bank at 340K+ borrowers; SBI pilot scaling; HDB/L&T entering pipeline	Revenue diversification
Month 12–18	\$4.7M+	Full pipeline activation + second-wave bank references	Growth acceleration

The typical pattern for proven founders re-entering Indian BFSI: first 2–3 banks close quickly (relationship-driven), then 5–10 banks follow within 12 months (reference-driven). Once Yes Bank and SBI are live references, the second wave accelerates significantly—these are the banks every other Indian FI watches.

Validation Metrics: What We Will Measure

Upon production launch in April/May 2026, Krim will track and report the following metrics to quantify the pre-execution validation advantage over post-hoc approaches:

Metric	What It Measures	Target
Pre-execution block rate	% of actions blocked by validators before execution that would have passed post-hoc filters	Reported from Month 1 of production
Compliance incident rate	Regulatory violations per 10,000 operations vs. industry baseline	Order-of-magnitude reduction vs. manual baseline
Validator-specific catch rate	Which of the 26 validators catch the most failures, by type and severity	Continuous calibration data
False positive rate	% of valid actions incorrectly blocked by validators	<2% target to ensure operational throughput
Audit trail completeness	% of actions with full epistemic justification logged	100% from Day 1

These metrics will provide the hard performance data that converts the epistemological thesis into measurable compliance infrastructure. Early production data from Yes Bank will be incorporated into the Seed raise materials.

9. Go-to-Market Strategy & Scaling Partnerships

Market Opportunity

Market	Size	Why It's Ready
India End-to-End Credit Ops	\$5B+	Debt collection (\$850M), loan servicing (\$1.5–2B), credit BPO (\$1–1.5B), collections tech (\$500M); growing 15–18%
US Credit Servicing	\$50B+	\$15B compliance tax, CFPB enforcement surge, delinquencies at 15-year highs
Global TAM	\$140B+	Insurance, healthcare, legal—every regulated domain needing validated AI

Go-to-Market: Relationship-Led, Then Reference-Driven

Krim's GTM strategy leverages the founder's established relationships across 50+ Indian banks and NBFCs—the most capital-efficient distribution channel available at pre-seed. The approach operates in two phases:

Phase 1: Relationship-driven (Months 1–6). Yes Bank and SBI represent the founder's direct relationships. These engagements were secured before any institutional capital, demonstrating that the founder's pedigree—48M borrowers, zero violations—is itself a distribution asset. HDB Financial Services and L&T Finance are in discovery through the same network. The first 2–3 bank deployments close through trust built over 7+ years of production relationships.

Phase 2: Reference-driven (Months 6–15). Once Yes Bank and SBI are live, the reference effect activates. In Indian BFSI, these are the institutions every other bank watches. The typical pattern: 5–10 banks follow within 12 months once marquee references are live. Production data and compliance metrics from Phase 1 deployments become the sales collateral for Phase 2.

Scaling Strategy: Ecosystem Partnerships for Accelerated Penetration

While relationship-driven GTM is capital-efficient, scaling beyond the founder's direct network requires strategic partnerships with organisations that have deep, existing penetration into the banking ecosystem. Krim's partnership strategy targets three categories:

1. Core Banking & Technology Integrators. India's banking technology stack is dominated by a small number of core banking system providers (Infosys Finacle, TCS BaNCS, Oracle FLEXCUBE) and system integrators (Wipro, Infosys, TCS) that maintain deep, long-standing relationships with every major bank. A technology partnership or integration with a core banking provider embeds KrimOS into the deployment pipeline that banks already trust and procure through. This is the highest-leverage partnership vector for scaling beyond door-to-door sales.

2. Industry Bodies & Regulatory Sandboxes. RBI’s regulatory sandbox framework and industry bodies like IDRBT (Institute for Development and Research in Banking Technology) and IBA (Indian Banks’ Association) offer pathways to credentialed endorsement. Participation in regulatory sandboxes provides both validation and visibility across the banking ecosystem. Academic collaboration with JNU/IIT already positions Krim for regulatory engagement through published research.

3. Investor & Advisor Networks. Pre-seed and Seed investors with BFSI portfolios bring not just capital but distribution. Investors with existing banking relationships—particularly those who have backed other BFSI infrastructure companies—provide warm introductions that bypass cold procurement cycles. Advisory board members with C-suite banking relationships further accelerate pipeline development.

Competitive Positioning in Partnerships

Notably, several potential partners have built or acquired their own competing capabilities. Yubi (which acquired Spocto and launched YuVerse.ai) and Perfios (which acquired CreditNirvana’s collections platform) are building their own agentic AI and collections intelligence layers. These are competitors, not partners. Krim’s partnership strategy explicitly targets organisations with deep ecosystem access but without competing AI ambitions in the credit operations space.

Indo-US Dual Market Strategy

	India	US
Role	Scale proving ground + early revenue	Revenue multiplier + valuation driver
Customers	SBI, Yes Bank, HDB, L&T Finance	Credit unions, auto lenders, mid-market banks
Advantage	500M+ customers, rapid deployment, founder’s bank network	\$50B market, 10× pricing, regulatory complexity = moat
Timeline	Production from April/May 2026	Seed-funded expansion: first US pilots Month 12–18

US market entry is a Seed-stage initiative, not a pre-seed goal. Louis Parkinson (COO, New York) provides US market access and enterprise relationships. Active US selling begins after Indian deployments are live and referenceable.

10. The Instrument: \$1.5M CCPS

Term	Detail
Instrument	Compulsory Convertible Preference Shares (CCPS)
Amount	\$1,500,000
Valuation (Pre-Money)	\$7.5M pre-money
Conversion Discount	20% discount to Seed price
Expected Seed	\$15–20M raise at \$100M+ valuation in 6–15 months
Return Profile	10–20× in 12–18 months; 25–50× over 3–5 years (including 20% discount)

Use of Funds

Allocation	Amount	Purpose
Engineering (45%)	\$675K	KrimOS runtime scaling, Karta agent library, VP Eng hire within 3–4 months
Compliance & IP (15%)	\$225K	SOC2 Type 1, RBI sandbox, provisional patent filings, trade secret documentation
GTM & Bank Deployments (25%)	\$375K	Yes Bank production, SBI pilot scaling, HDB/L&T pipeline development
Working Capital (15%)	\$225K	Operational continuity, legal, administration

Capital efficiency note: With ~\$1M ARR from Yes Bank alone (₹70L MRR from Month 1), Krim reaches partial self-funding within months of production deployment. The pre-seed capital is bridge-to-revenue, not runway-to-next-raise—a fundamentally healthier capital structure than most pre-seed investments.

11. Risks & Mitigations

Risk	Severity	Mitigation
Execution speed	Low	Same team: 200M+ calls, 50+ banks, zero violations at Saarthi. Working GTM product in bank testing and demo.
Bank procurement timelines	Medium	Yes Bank onboarding active; SBI LOI expected before trial. Founder has navigated PSB procurement dozens of times. Relationship layer—the slowest part—is already cleared.
Product readiness for full-stack operations	Medium	Working GTM product is in active iteration. Full Safe Superintelligence OS requires ~3 months additional development, shaped by bank team feedback. This is deliberate: building with production requirements ensures real-world fit.
Banks build internally	Low	Requires AI + epistemology + regulatory + scale expertise simultaneously. 3–5 year gap minimum. No bank has this interdisciplinary team.
Competitive entry from large players	Low	26 validators require rare interdisciplinary expertise. IP strategy (patents + trade secrets) in place. 18–24 month head start. Compounding domain embedding creates structural gap that widens with each deployment.
Founder bandwidth	Medium	Triple role (CEO/CTO/CAIO) is managed via VP Eng hire in Month 3–4. Two co-founders with distinct domains.
US market entry	Low	Deferred to Seed stage. Louis in NY provides access. India deployments create referenceable proof first.
Partnership ecosystem limitations	Medium	Key ecosystem players (Yubi/YuVerse, Perfios) are building competing capabilities. Krim targets non-competing partners: core banking integrators, regulatory sandboxes, and investor/advisor networks with banking relationships.

12. The Vision: Beyond Credit Operations

Phase	Timeline	Domain	Revenue Potential
Phase 1	Year 1–2	Post-disbursal credit operations (autonomous collections, servicing)	\$10M–\$50M ARR
Phase 2	Year 2–3	Full autonomous banking operations + insurance claims (replaces current SaaS platforms)	\$50M–\$150M ARR
Phase 3	Year 3–5	Cross-industry safe superintelligence runtime: healthcare, legal, any regulated domain	\$150M+ ARR

The end state: autonomous banking operations powered by safe superintelligence infrastructure. KrimOS—powered by Kendra, the intelligence runtime where every AI action is validated before execution—becomes the platform that autonomous banking is built on, the way HTTPS became the standard for secure web traffic. As Epistemological AI matures, the same infrastructure extends to every regulated industry where AI actions have consequences.

The path from credit operations to autonomous banking to cross-industry runtime follows a proven expansion pattern: dominate a single vertical (post-disbursal credit ops), establish regulatory credibility and production data, expand across the full banking operations lifecycle (the transition from SaaS platforms to autonomous operations), then extend into adjacent regulated domains. Insurance claims processing, healthcare billing, and legal automation all share the same structural requirement: every AI action must be validated before execution. The 26 validators are domain-agnostic in principle; only the calibration data is domain-specific.

What Must Be True for Multi-Billion Outcome

1. AI maturity reaches the point where autonomous banking operations outperform human-middleware-based SaaS platforms on cost, speed, compliance, and quality—accelerating the replacement cycle.
2. Krim becomes the compliance standard in BFSI AI—regulators endorse or indirectly validate the epistemic framework as the required safety layer for autonomous operations.
3. Large banks mandate pre-execution epistemic validation for AI agent deployment, making Kendra the default intelligence runtime for autonomous banking.
4. Current SaaS banking platforms (collections, servicing, operations) are displaced by Krim’s autonomous operations architecture—not digitising human workflows but replacing them entirely.
5. Krim expands beyond banking into insurance, healthcare, and legal—Kendra becomes the horizontal safe superintelligence runtime for every regulated industry.



Sovereign Safe Superintelligence for Autonomous Banking

“You can’t make banking autonomous by adding AI to broken workflows. You need a new runtime—one where every action is validated before execution. That’s what we’re building.”

— Shri Vishwa Nath Jha, Founder & CEO/CTO/CAIO

C O N T A C T

Shri Vishwa Nath Jha | nath@krim.ai | +1 (510) 345-5686 · +91 8618064285 | www.krim.ai

Krim AI Inc. | New York · Bangalore · Patna